JASMINE

Data Stream Processing

Simone Mancini Francesco Ottaviano Andrea Silvi

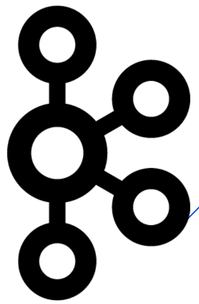

JASMINE

Design Choices

Frameworks

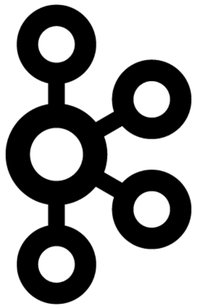

Programming Languages

Java

JavaScript

Store

redis

Queries

Query 1 - DAG

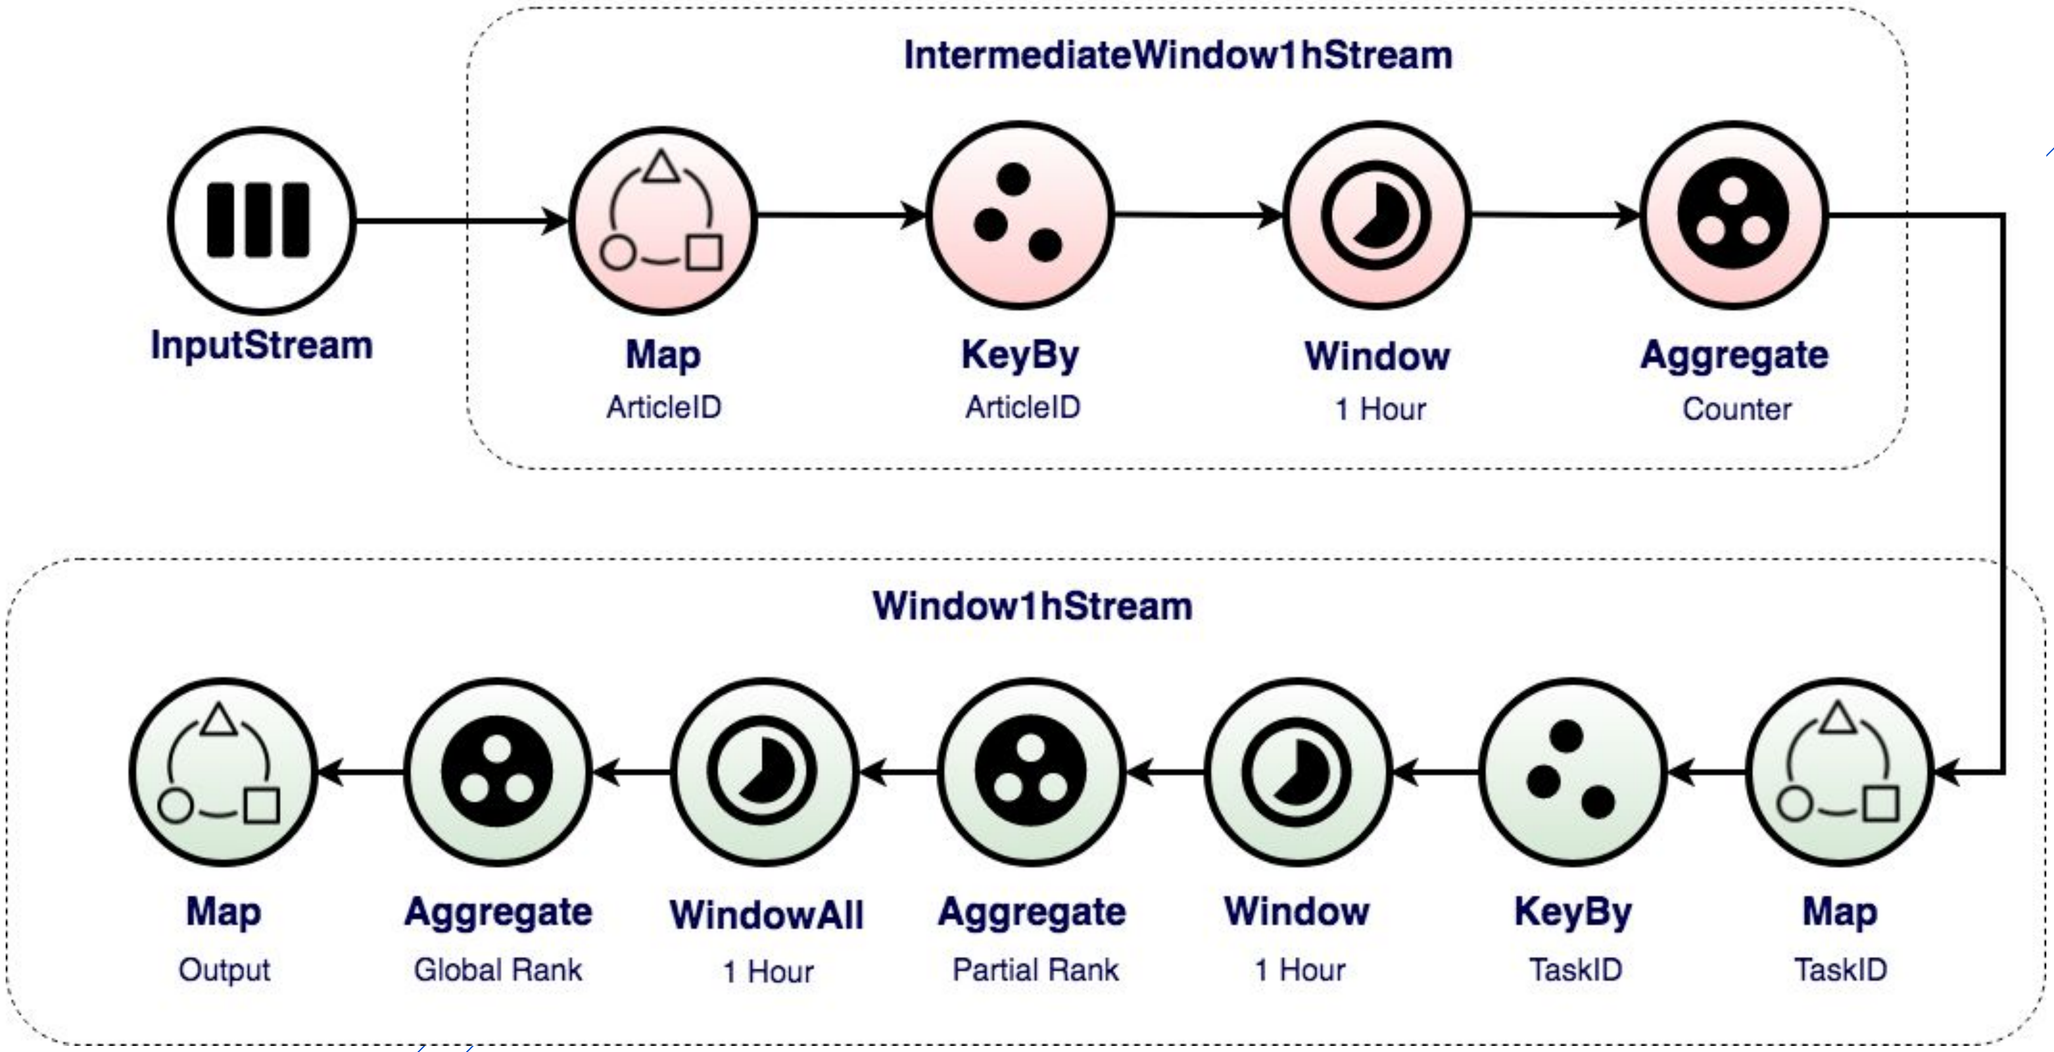
IntermediateWindow1hStream
InputStream
Map
ArticleID
KeyBy
ArticleID
Window
1 Hour
Aggregate
Counter
Window1hStream
Map
Output
Aggregate
Global Rank
WindowAll
1 Hour
Aggregate
Partial Rank
Window
1 Hour
KeyBy
TaskID
Map
TaskID

JASMINE

Query 1 - Chaining

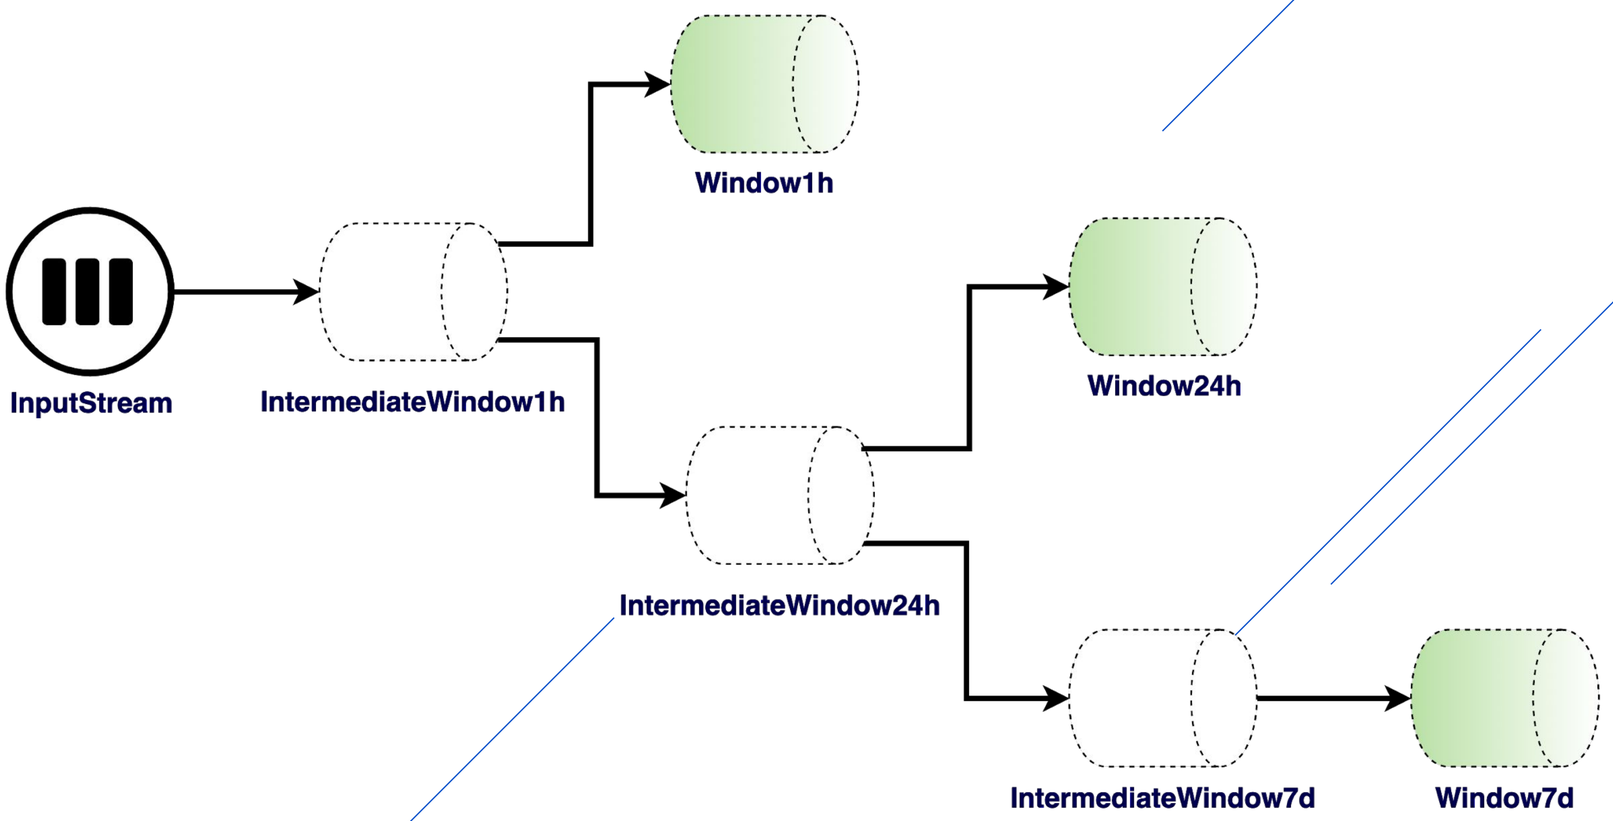

JASMINE

Query 2 - DAG

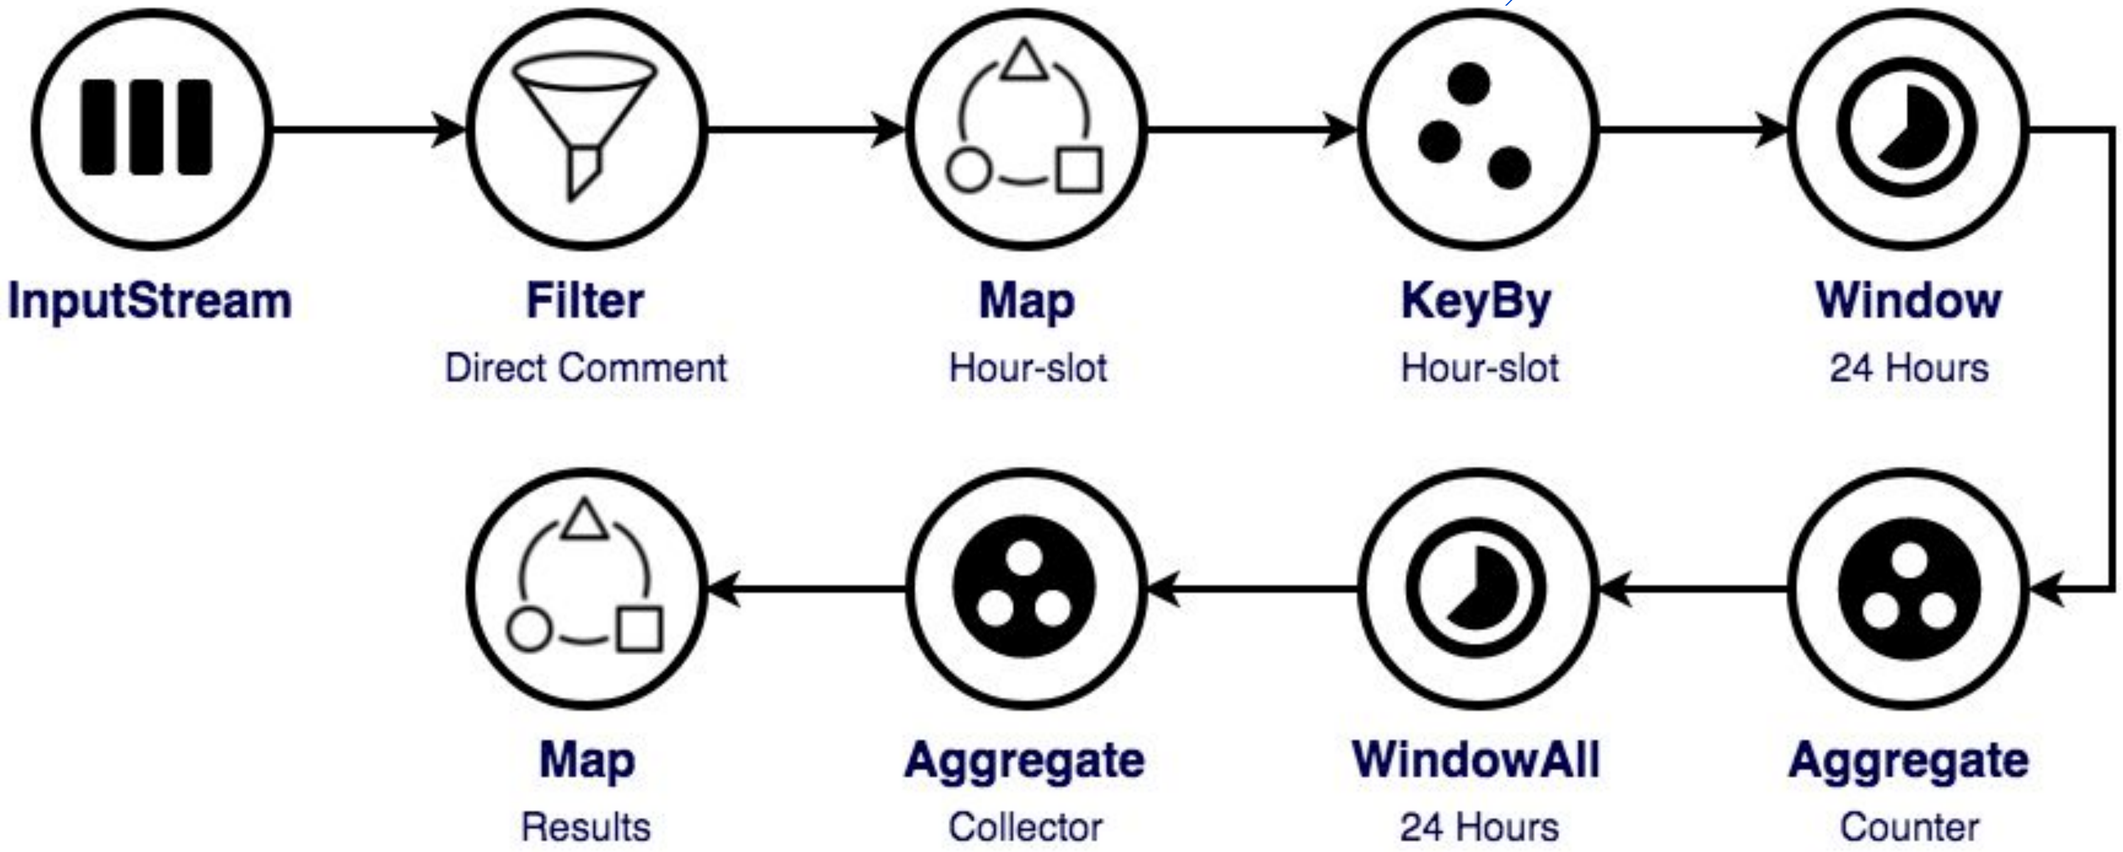

JASMINE

Query 2 - Chaining

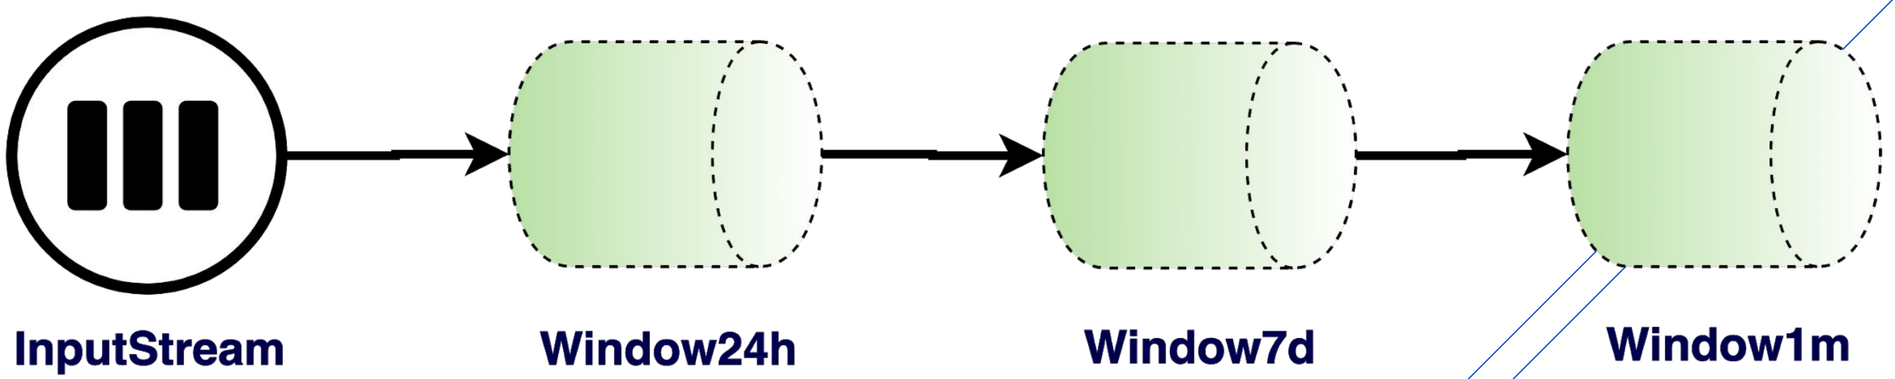

JASMINE

Query 3 - DAG

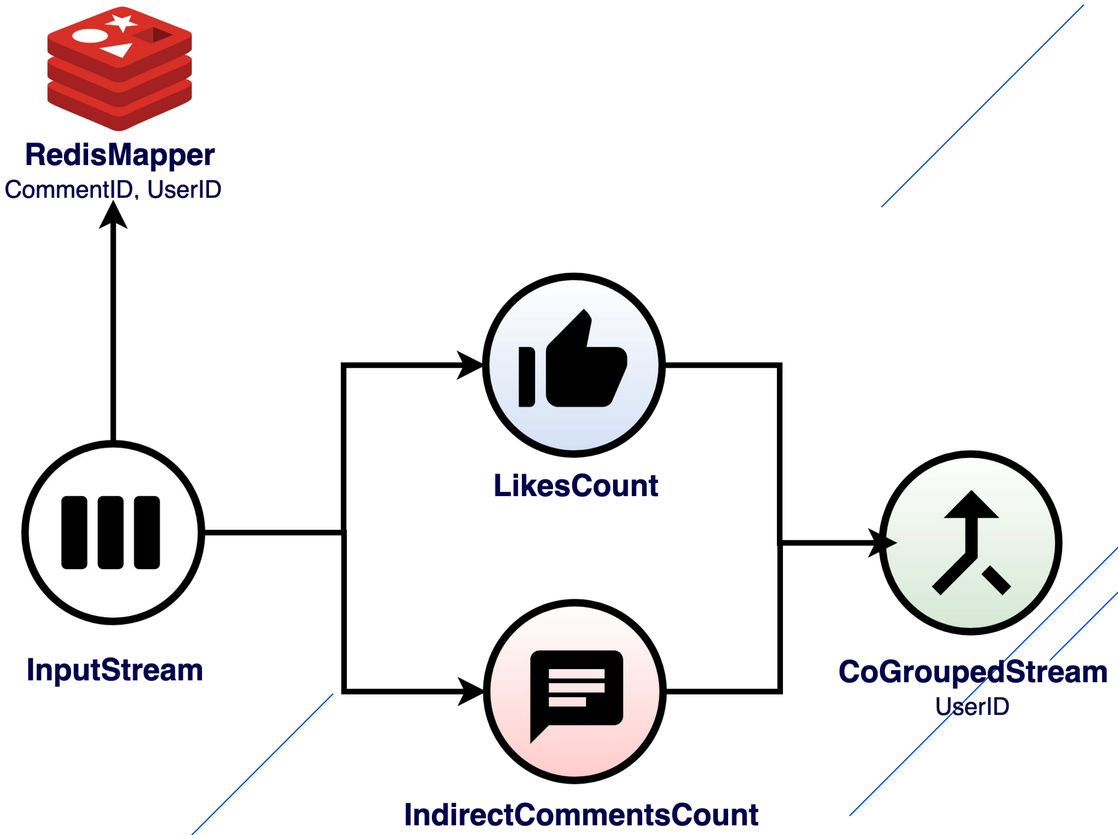

JASMINE

Query 3 - DAG

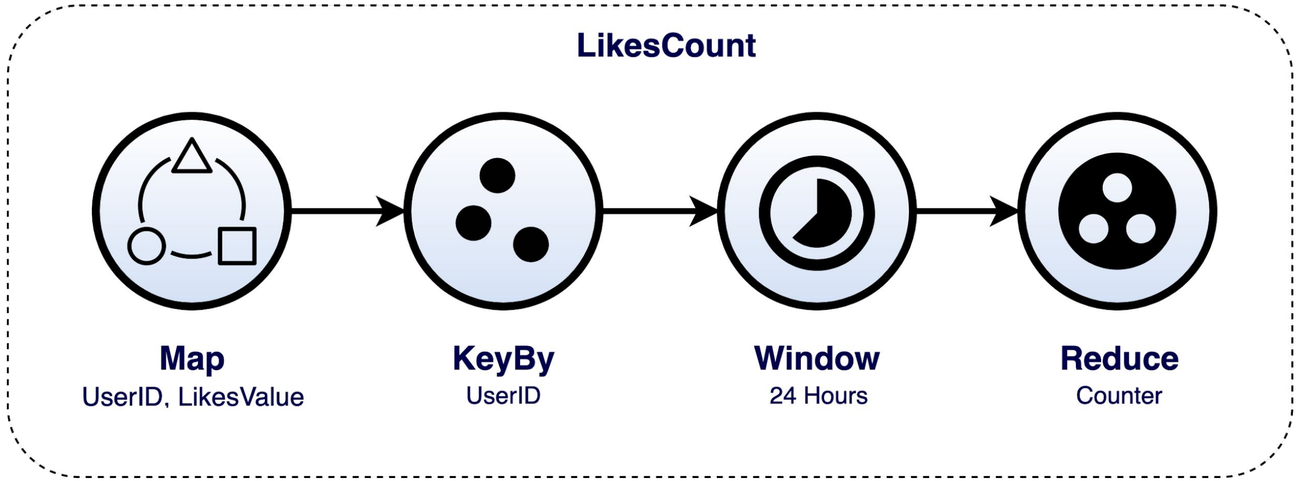

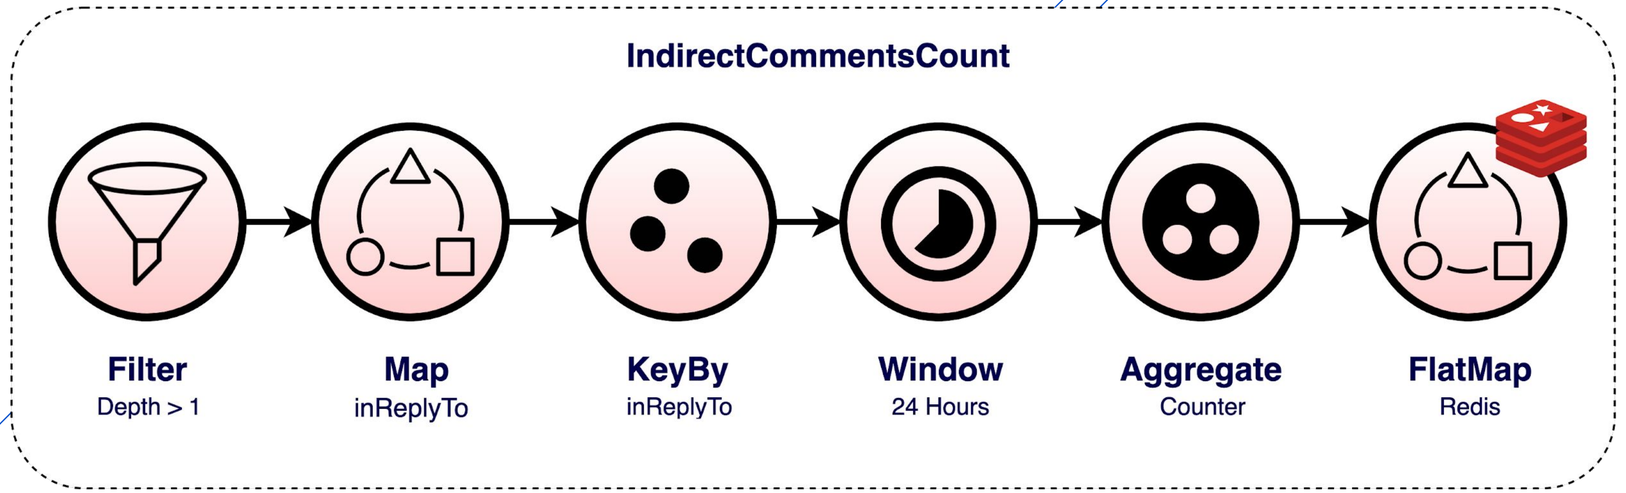

JASMINE

Query 3 - DAG

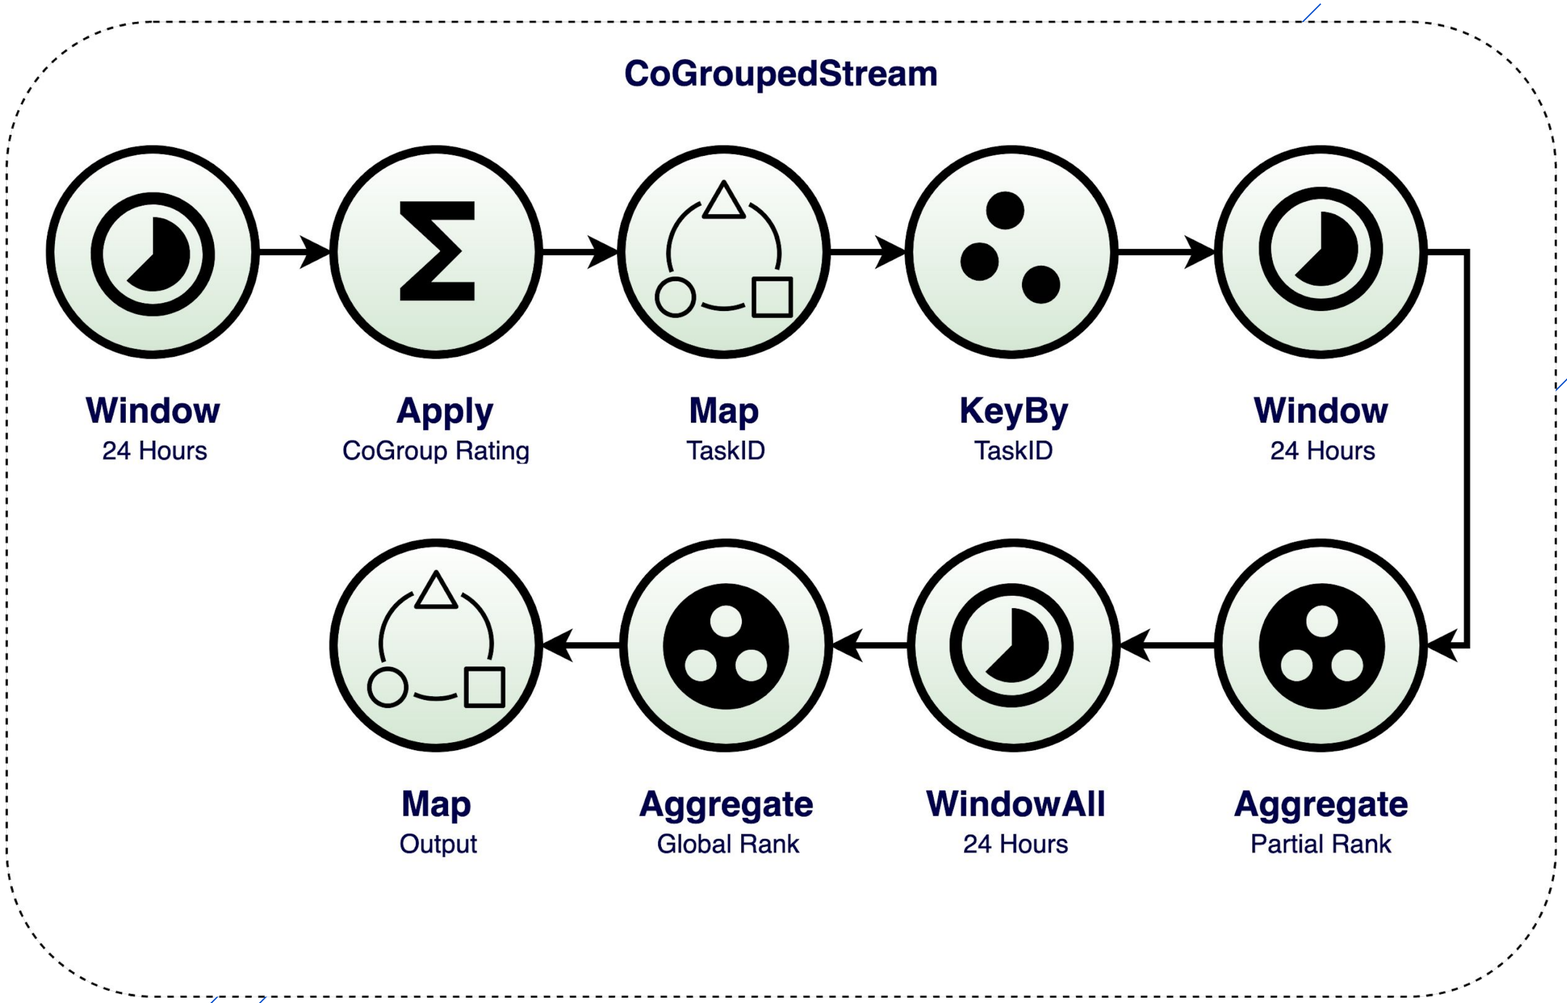

JASMINE

Query 3 - Chaining

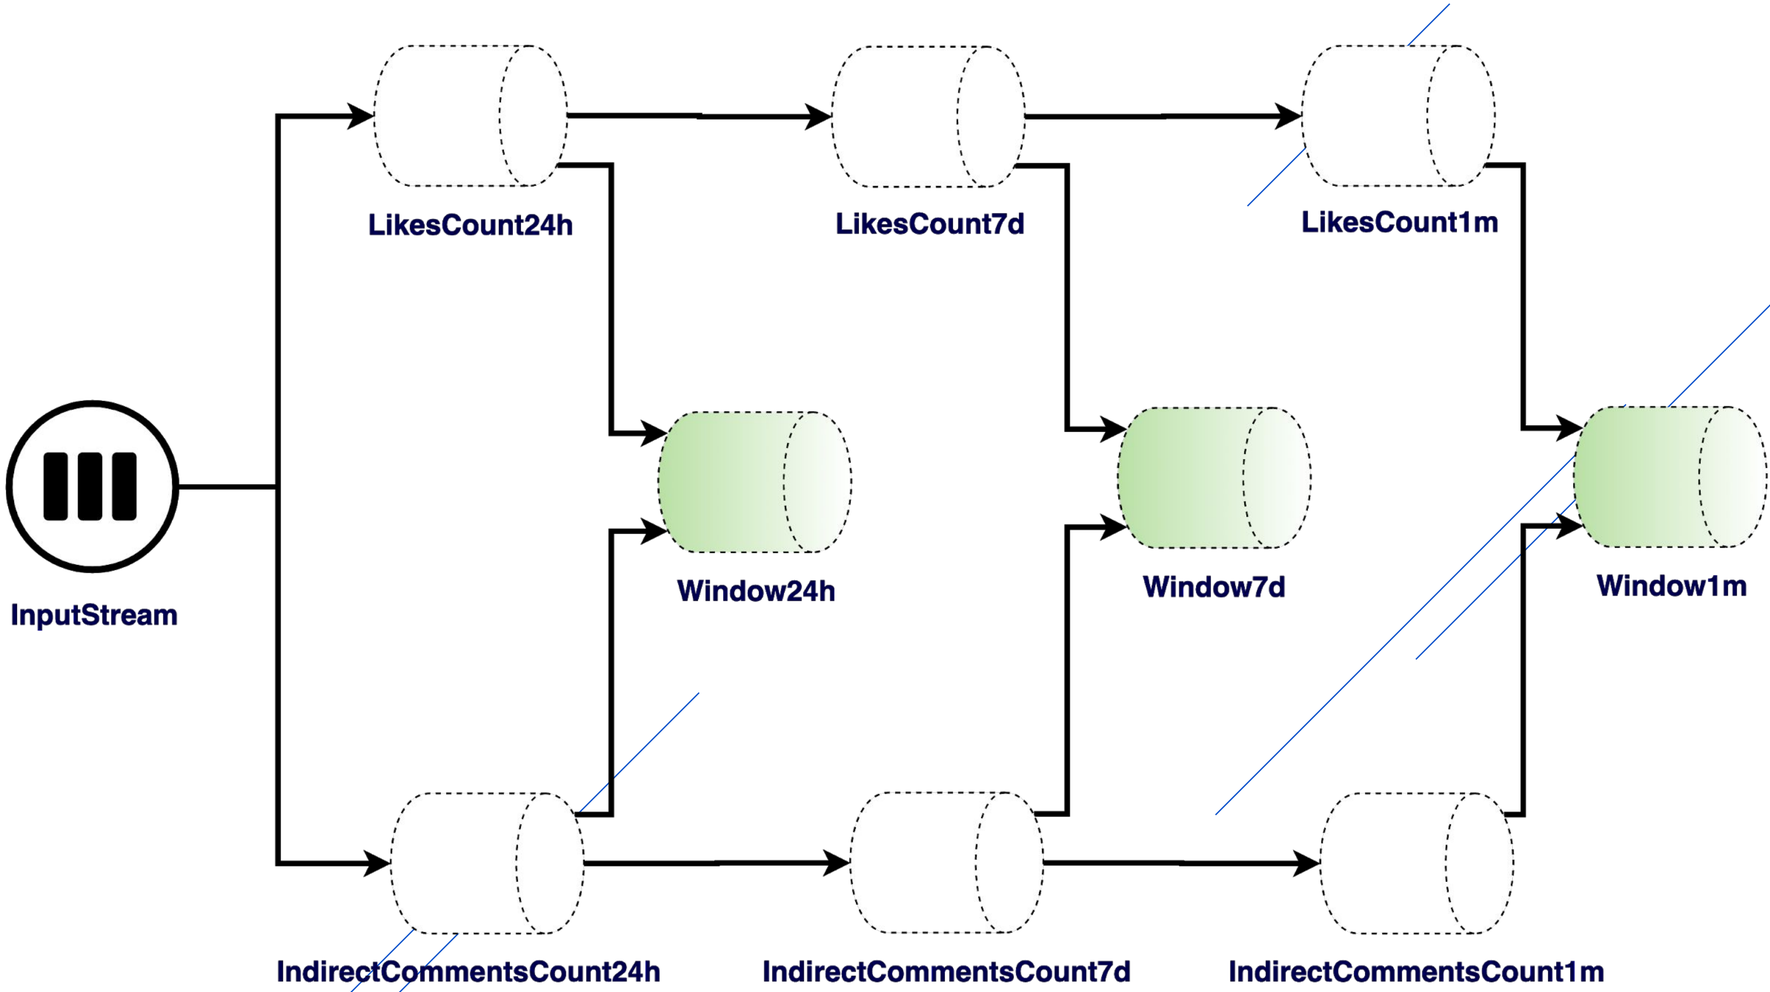

Kafka Streams

JASMINE

KAFKA STREAMS

- Different approach: *Kstreams* and *Ktable*
- Different ranking solution
- Parallelism
- *Suppress* in chaining mode limitations
- Custom Serde

Architecture

JASMINE

Architecture

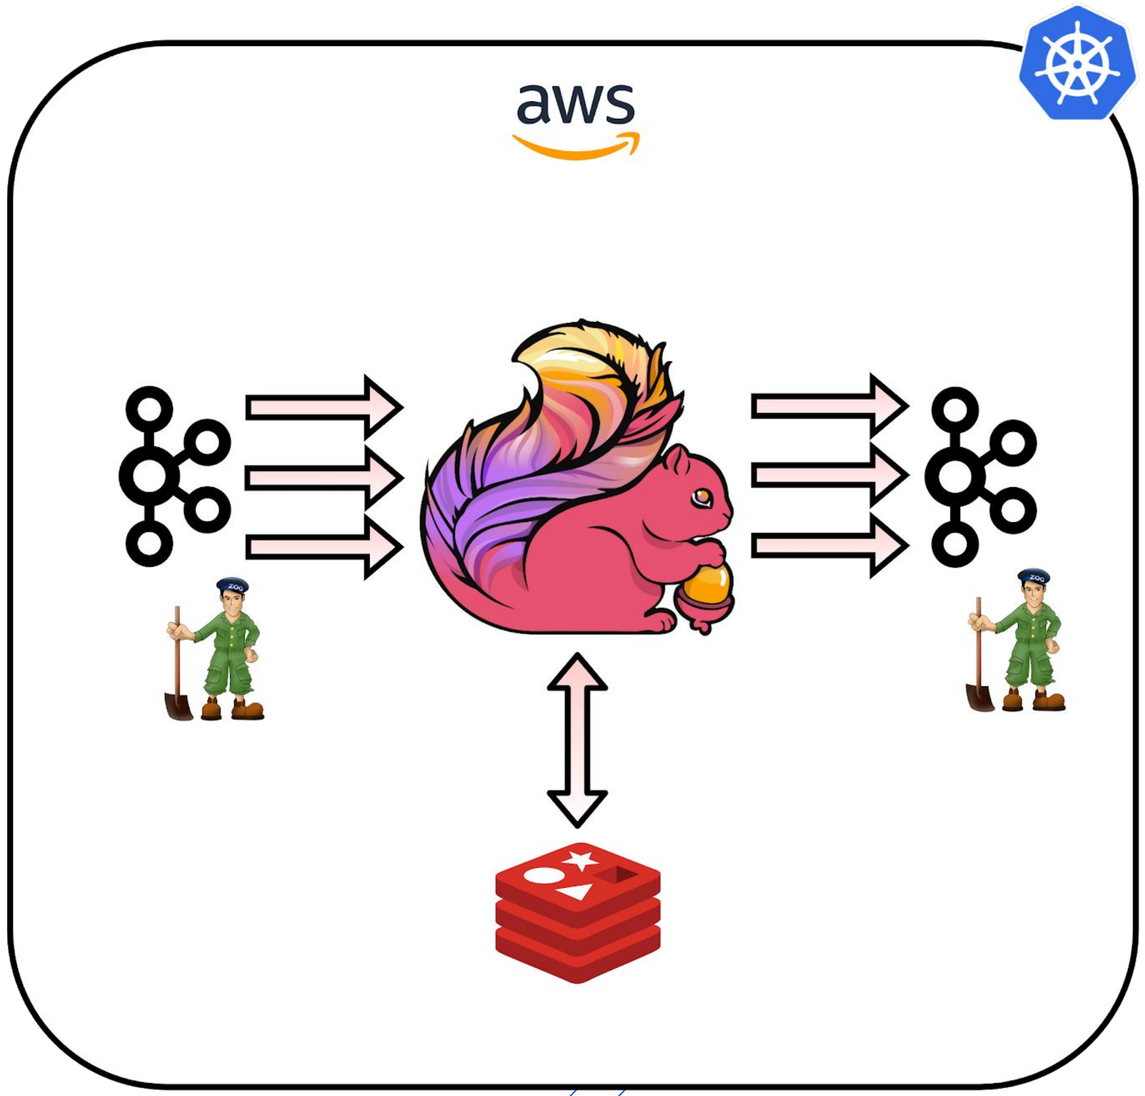

- *m4.large EC2 instances*
  - 2,4 GHz Intel Xeon E5-2676 v3 Processor
  - 2 vCPU, 8GiB Mem
- *Kubernetes*
  - *Kafka*: 1 node per machine, 3 nodes with replication factor 3 on topics.
  - *Zookeeper*: 1 node per machine, 3 nodes.
  - *Flink*: 1 job manager, 8 task manager.

JASMINE

Simulator

- Emulates data stream processing compressing time
- *NodeJS*
- *CreateDate* as tuple creation time

Evaluations

Apache Flink - Throughput

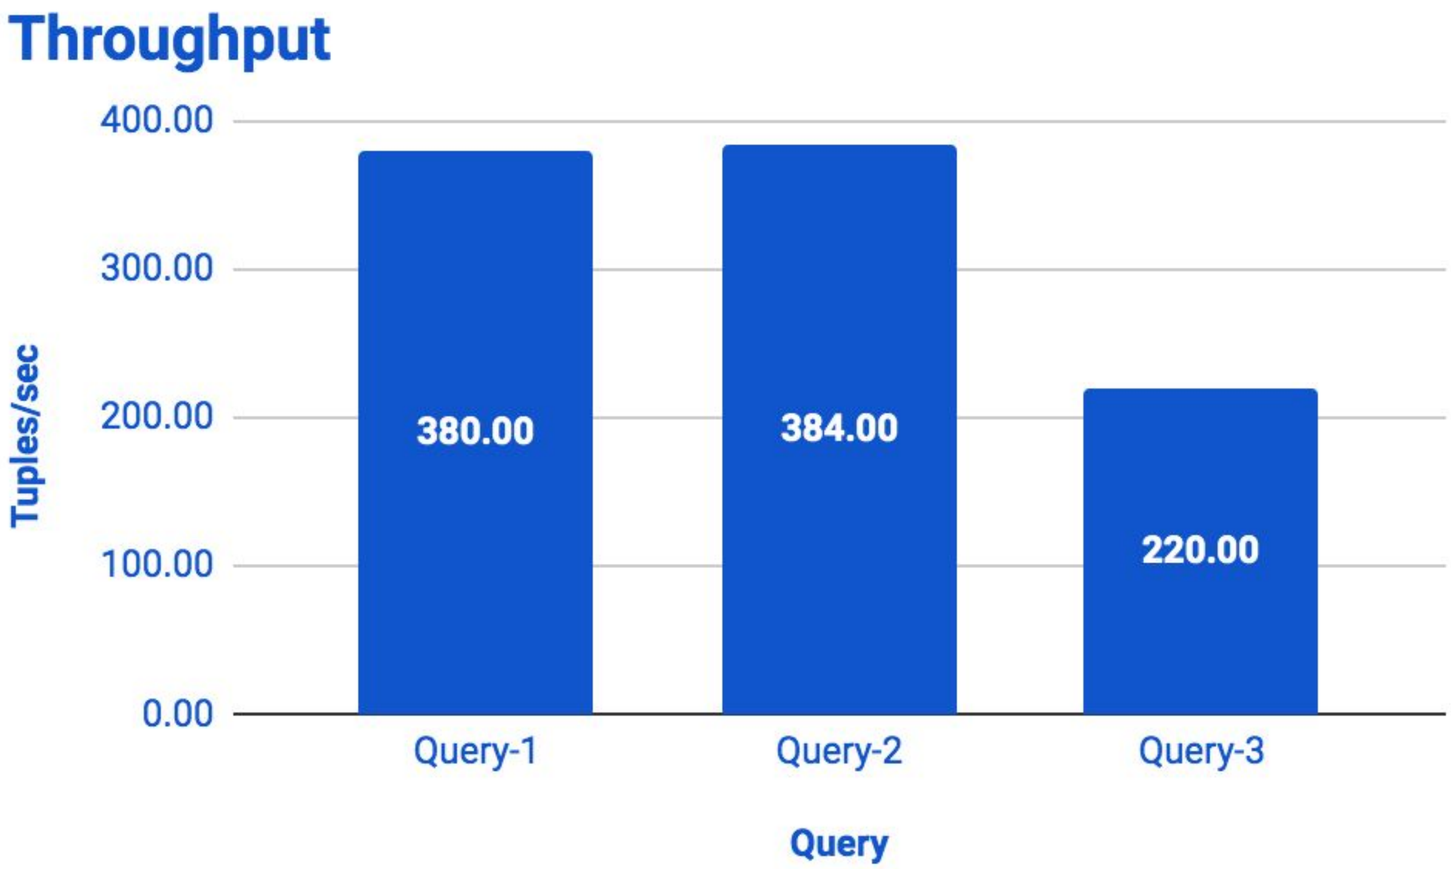

JASMINE

Apache Flink - Latency

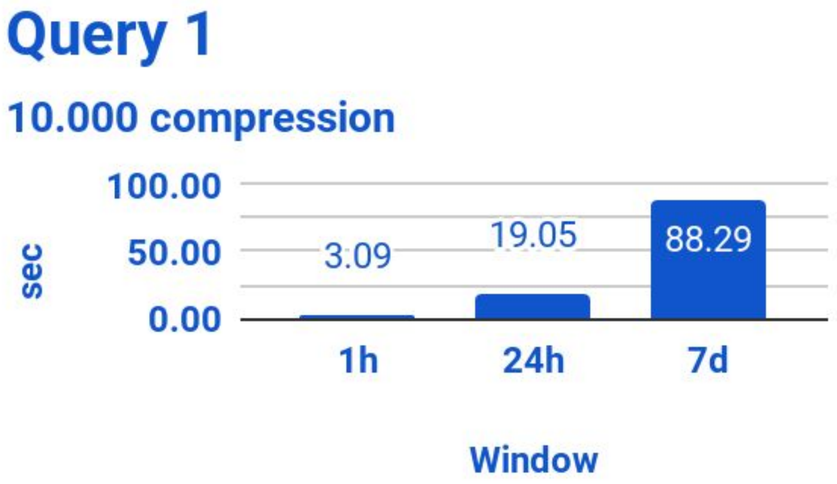

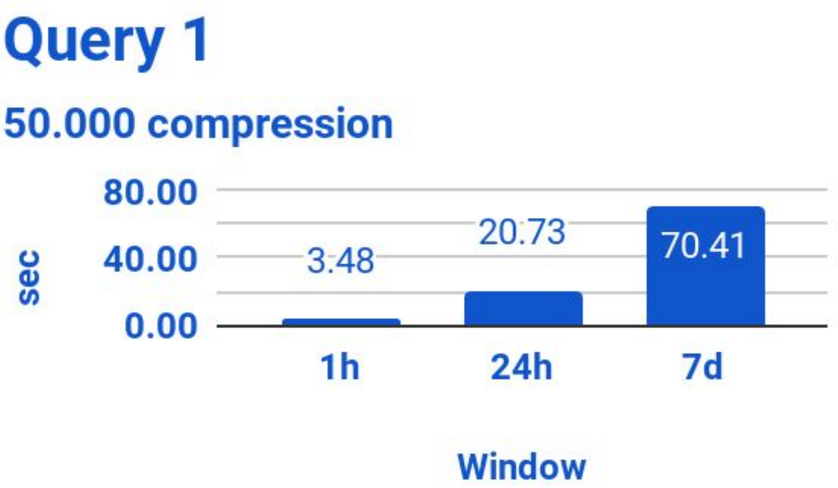

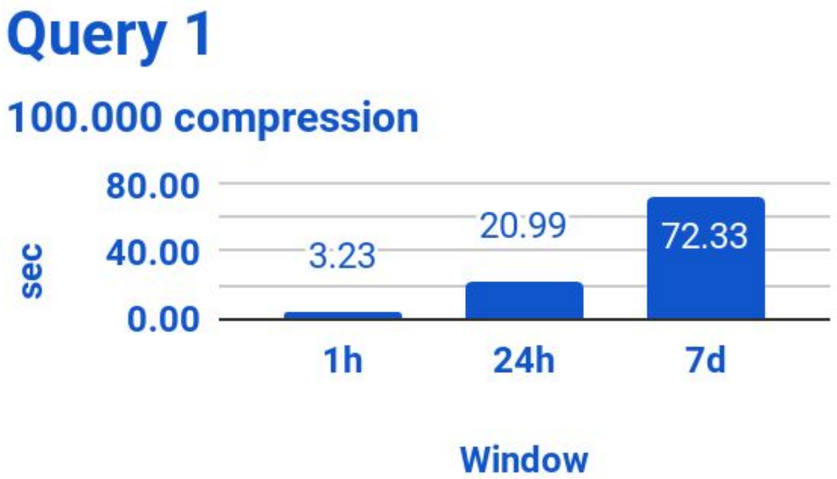

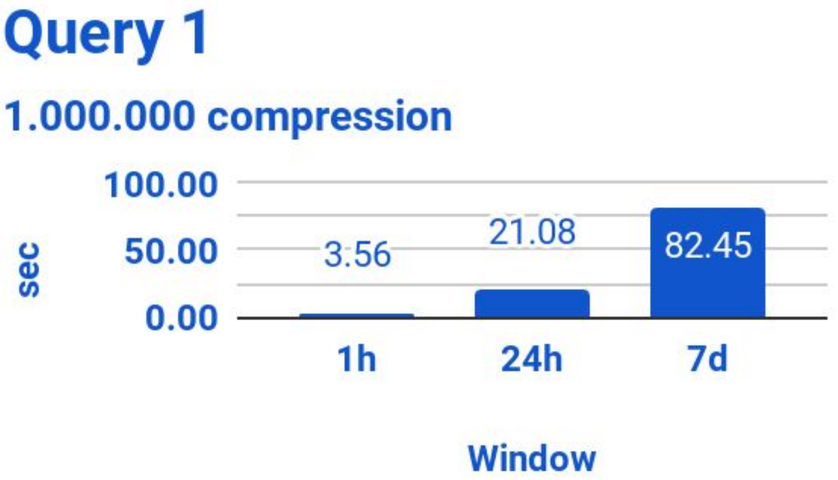

JASMINE

Apache Flink - Latency

Query 1

Windows Comparison by comprssion

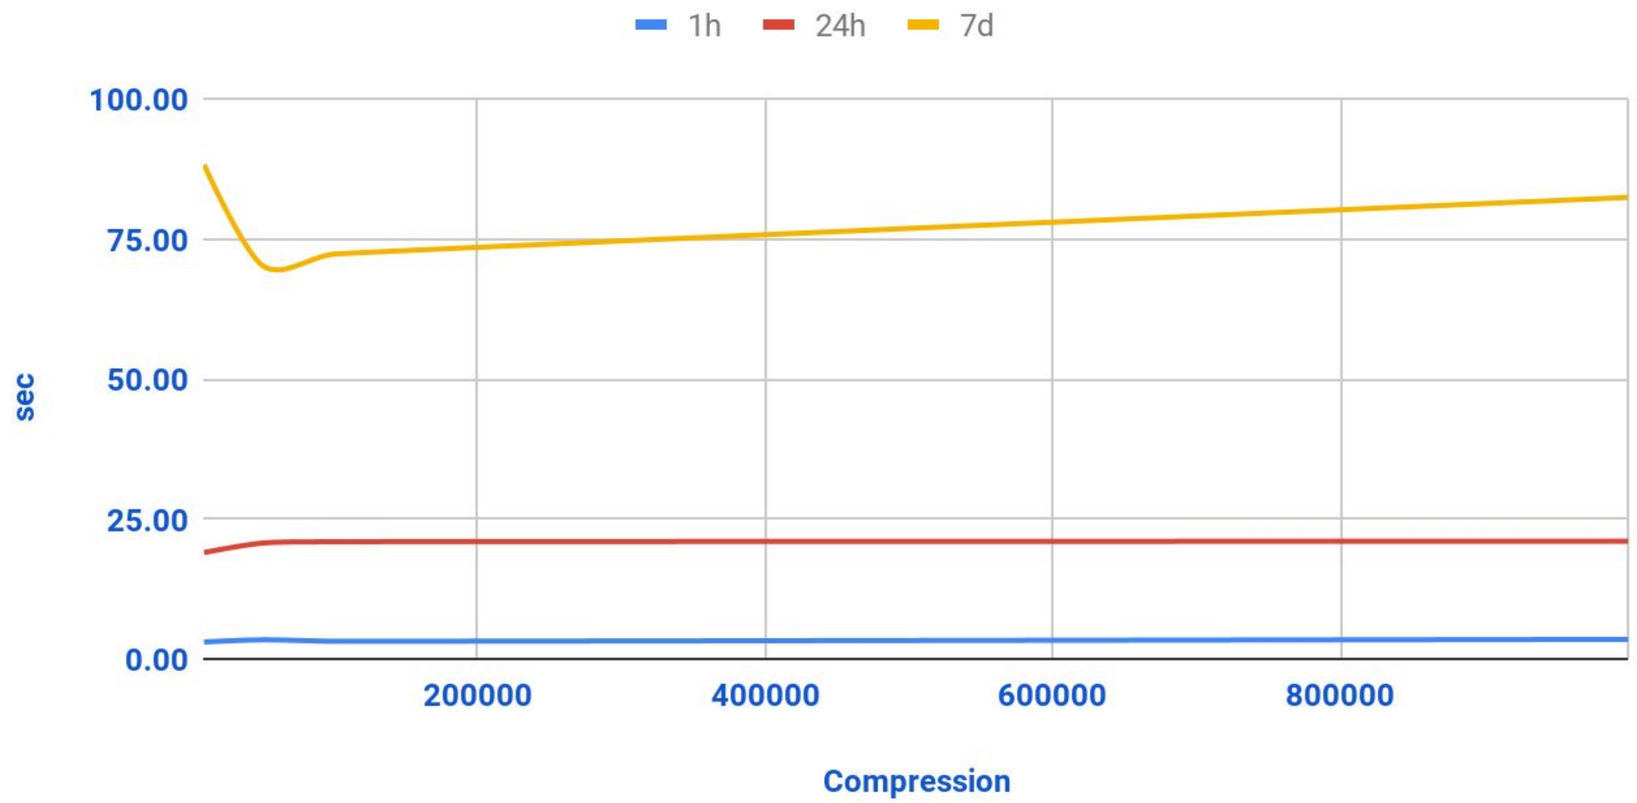

JASMINE

Apache Flink - Latency

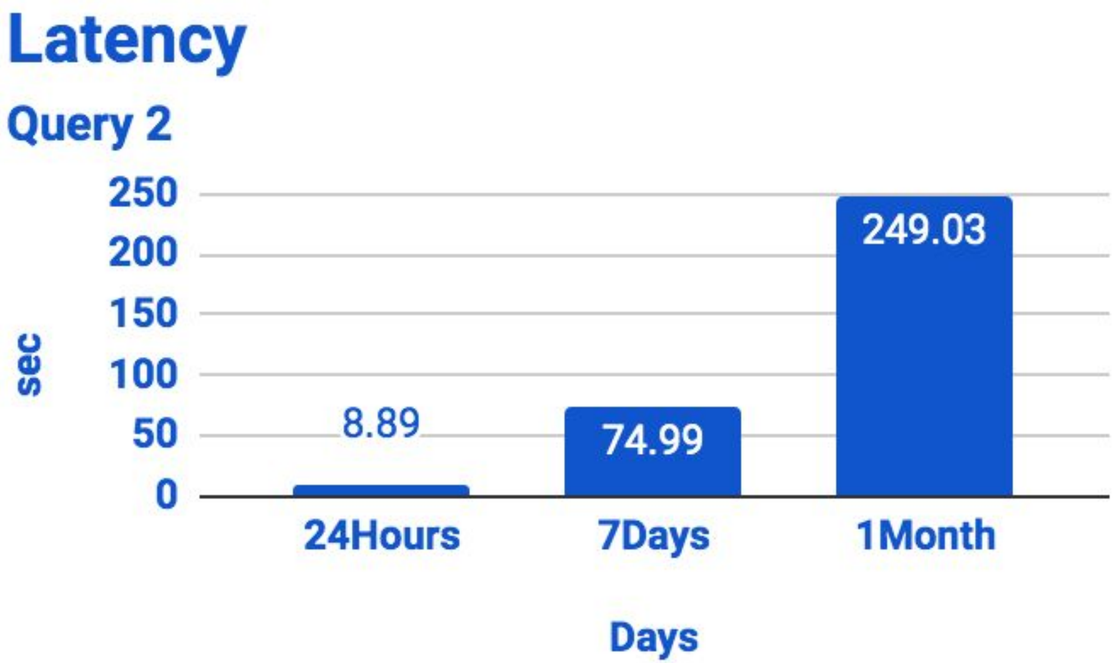

KafkaStreams - Throughput

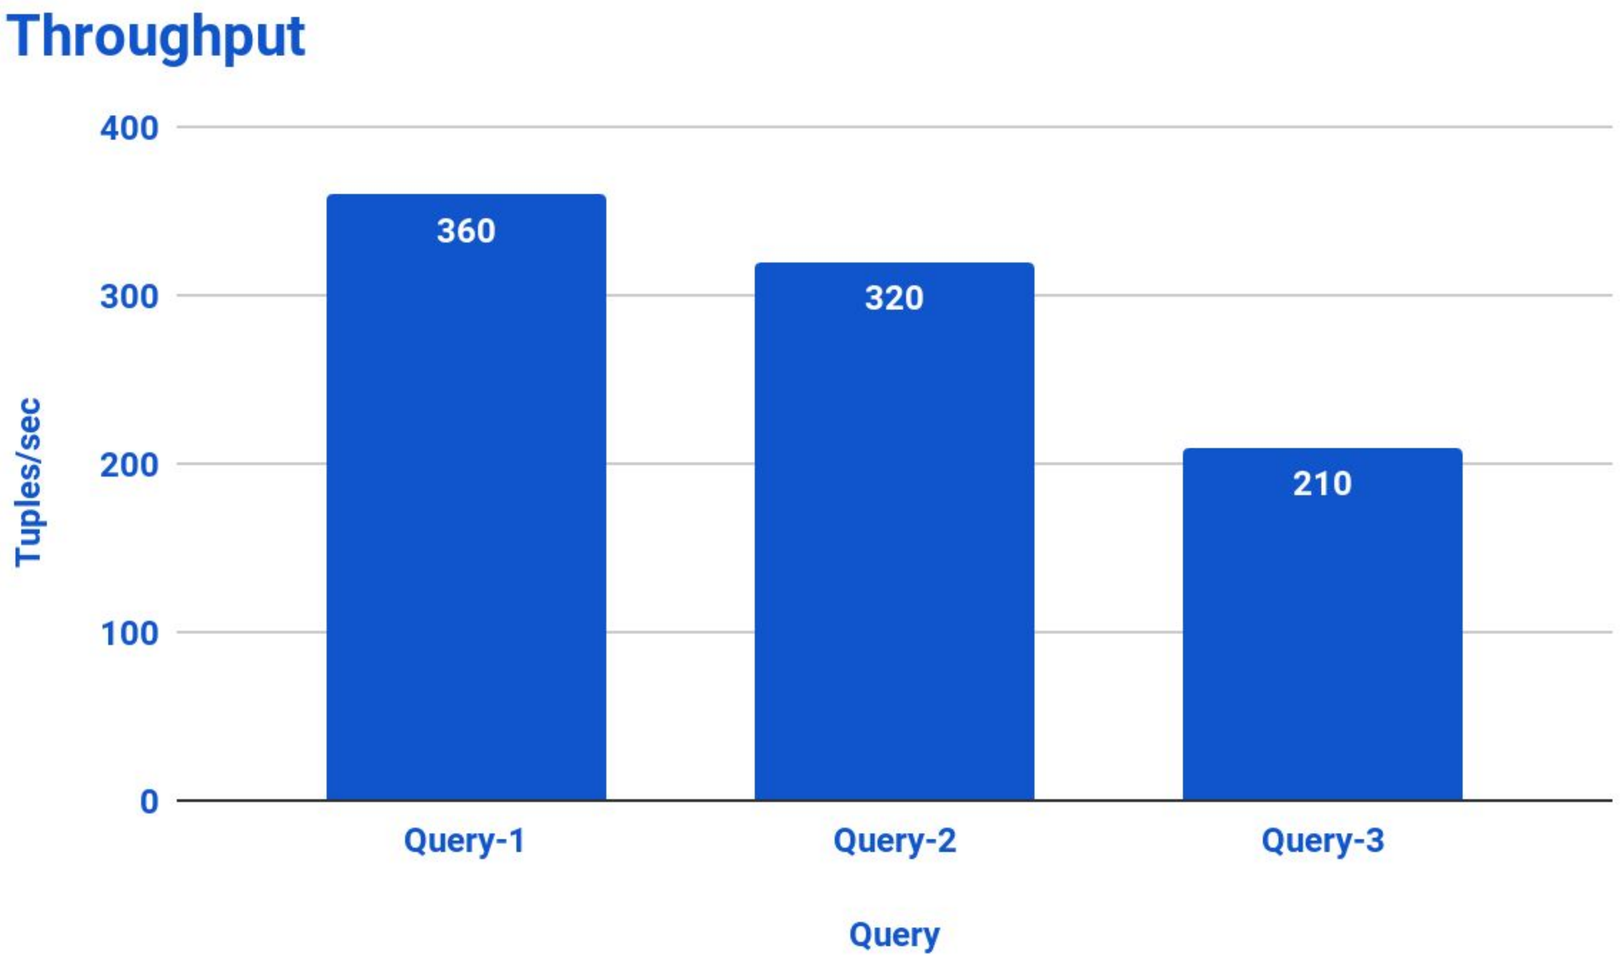

JASMINE

KafkaStreams - Latency

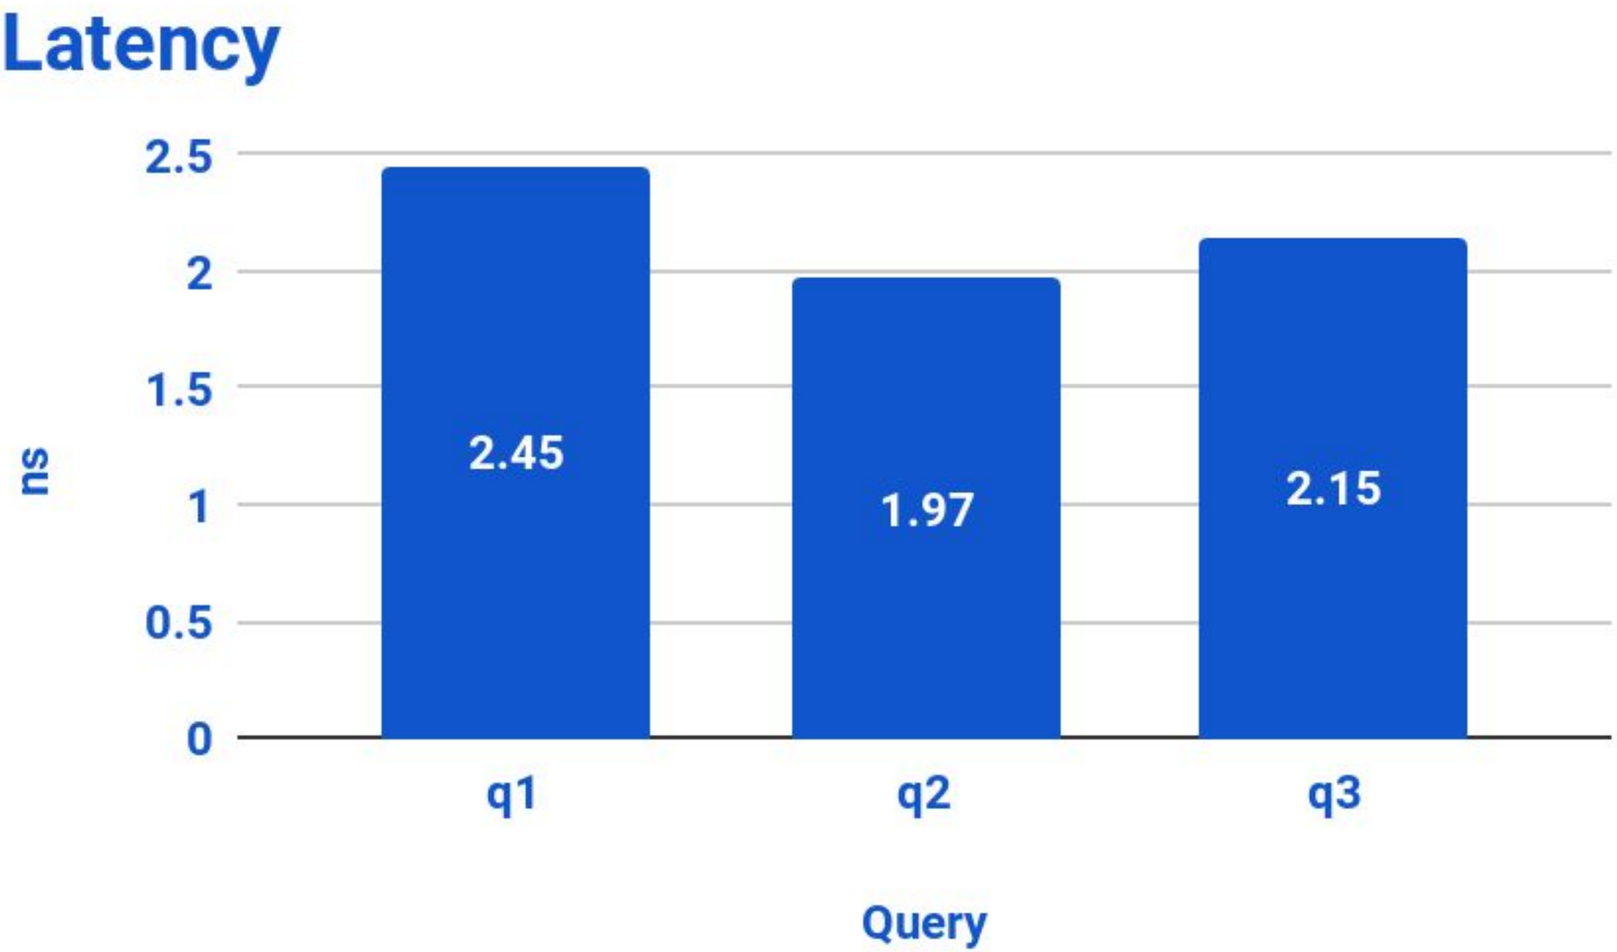

JASMINE

Apache Flink - Cluster evaluation

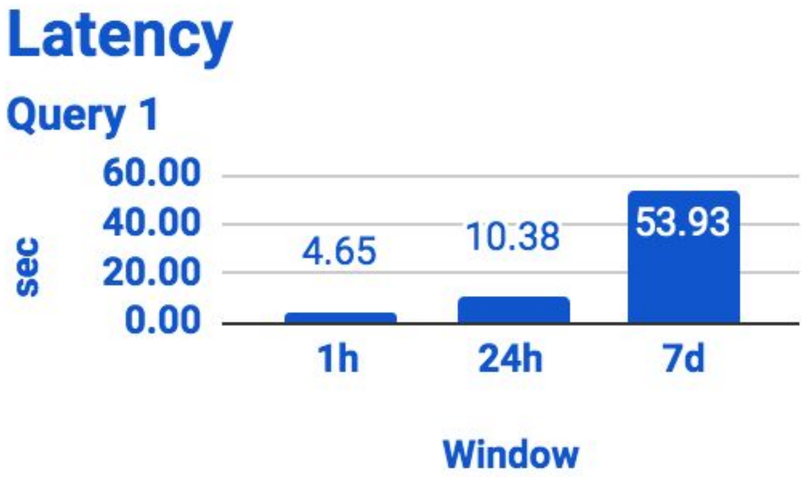

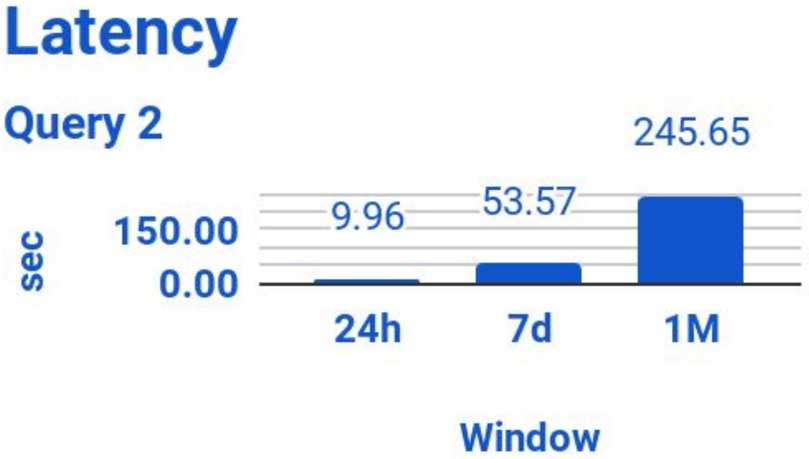

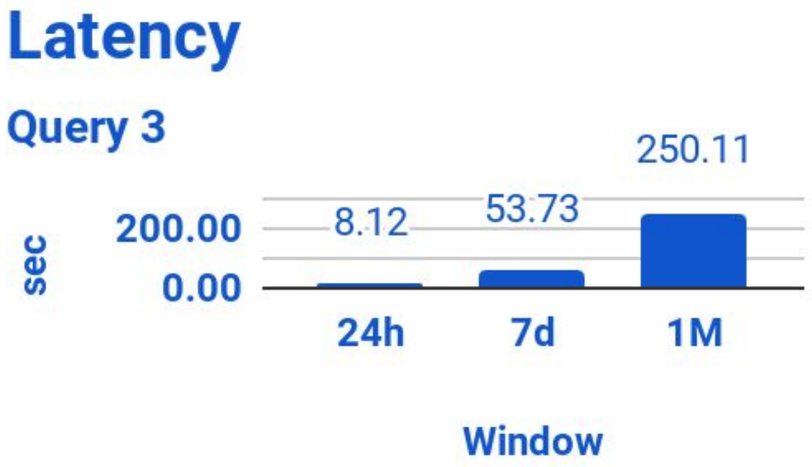

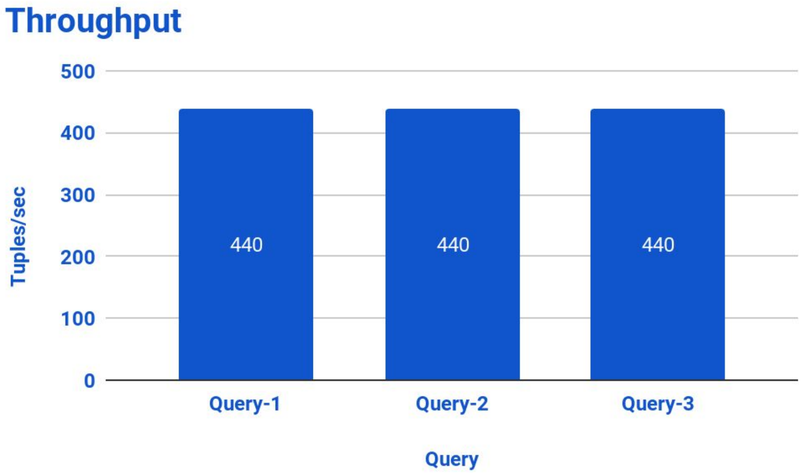

JASMINE

Thank You!